

John A. Mercer

AI Engineer

Contact Information

Email: john.mercer.ai@example.com

Phone: +1 (555) 482-1930

Location: San Francisco, CA, USA

LinkedIn: linkedin.com/in/johnmercer-ai

GitHub: github.com/jmercer-ai

Professional Summary

AI Engineer with 6+ years of experience designing, deploying, and optimizing machine learning and deep learning systems.

Specialized in Retrieval-Augmented Generation (RAG), LLM fine-tuning, vector databases, and scalable ML infrastructure.

Strong background in MLOps, model evaluation, and data engineering for production AI solutions.

Education

Master of Science in Artificial Intelligence – Stanford University (2017–2019)

Bachelor of Science in Computer Engineering – University of Texas at Austin (2013–2017)

Technical Skills

- Programming: Python, JavaScript/Node.js, Go, Bash
- ML & DL: PyTorch, TensorFlow, HuggingFace Transformers, LangChain, LangGraph
- LLM Ops: RAG pipelines, vector databases (Pinecone, Weaviate, FAISS), embedding optimization

- Cloud & DevOps: AWS, GCP, Docker, Kubernetes, CI/CD
- Data Engineering: Spark, Airflow, SQL, Parquet, data pipelines
- Tools: MLflow, Weights & Biases, Ray, DVC

Professional Experience

Senior AI Engineer – NovaMind Technologies (2022–Present)

- Designed and deployed a production-grade RAG pipeline with Pinecone and LlamaIndex handling 40M+ documents.
- Developed modular LLM agents for document extraction, entity recognition, and question answering.
- Implemented scalable fine-tuning workflows using HuggingFace Accelerate and LoRA adapters.
- Optimized inference latency by 35% via model quantization and caching strategies.

Machine Learning Engineer – DeepVision Labs (2019–2022)

- Built multi-stage classification and segmentation pipelines using PyTorch and timm.
- Led migration of ML infrastructure from on-prem to AWS EKS, reducing compute cost by 28%.
- Developed distributed training workflows with PyTorch Lightning and Ray.
- Implemented ML observability dashboards for model drift and performance monitoring.

Projects

Autonomous RAG Orchestrator

A modular framework that uses LangGraph to coordinate agents for retrieval, reasoning, summarization, and evaluation.

Technologies: OpenAI API, Pinecone, LangChain, FastAPI

LLM-Based Resume Evaluator

An AI system capable of parsing, embedding, and rating resumes using hybrid search and domain-tuned embeddings.

Certifications

- AWS Certified Machine Learning Engineer – Specialty
- Google Professional Cloud Architect
- Stanford NLP Certificate

Languages

English (Native), Spanish (Intermediate), German (Basic)